

# Probabilistic Robotics Course

## Finding Neighbors

**G. Grisetti**

`{grisetti}@diag.uniroma1.it`

Department of Computer, Control, and Management Engineering  
Sapienza University of Rome

# Neighbor Search

Given:

- a collection of vectors  $\mathcal{P} = \{\mathbf{p}_n\}_{n=1:N}$   $\mathbf{p}_n \in \mathbb{R}^k$
- a query vector  $\mathbf{p}_q \in \mathbb{R}^k$
- a distance metric  $d(\mathbf{p}_n, \mathbf{p}_q) \in \mathbb{R}^+$

Find either:

- the point in the collection that is the **closest** to the query, according to the metric

$$\mathbf{p}_i = \operatorname{argmin}_{\mathbf{p}_n \in \mathcal{P}} d(\mathbf{p}_q, \mathbf{p}_n)$$

- the points in the collection whose distance from the query is smaller than a value  $\epsilon$

$$\mathcal{P}' = \{\mathbf{p}_i \in \mathcal{P}, d(\mathbf{p}_q, \mathbf{p}_i) < \epsilon\}$$

# Distance Metrics

Examples:

- Squared Norm

$$\|\mathbf{p}_i - \mathbf{p}_j\|^2 = (\mathbf{p}_i - \mathbf{p}_j)^T (\mathbf{p}_i - \mathbf{p}_j)$$

- Omega Norm

this should look familiar

$$\|\mathbf{p}_i - \mathbf{p}_j\|_{\Omega}^2 = (\mathbf{p}_i - \mathbf{p}_j)^T \Omega (\mathbf{p}_i - \mathbf{p}_j)$$

- Hamming distance (for binary descriptors)

*Integer valued distance between two bit strings having the same dimension. Its value is the number of different bits.*

example:

`hamming(100100,101000)= 2`

# Trivial Approach

## Brute Force:

compute the distance metric between the query point and *each* of the points in the collection and update the minimum.

*Complexity:*  $O(N * \text{cost\_distance\_metric})$

If we need to perform many queries, this results in unacceptable delays.

**Idea:** *use auxiliary search structures.*

# Distance Map


If

- the dimension of the vectors is small ( $< 3$ )
- they are spread in a relatively small region of the space

we can *pre-compute* a **grid lookup table**.

Each cell of the grid contains:

- the distance from the closest point.
- the identity of the closest point.



# Distance Map


If

- the dimension of the vectors is small ( $< 3$ )
- they are spread in a relatively small region of the space

we can *pre-compute* a **grid lookup table**.

Each cell of the grid contains:

- the distance from the closest point (gray value)
- the identity of the closest point (orange line)




# Computing the Distance Map

Modified Version of *Dijkstra algorithm*.

Each grid cell  $C$  stores:

- $d$ : distance from nearest point
  - $\text{parent}$ : pointer to the nearest point
- We initialize the grid cells which corresponds to the points in  $\mathcal{P}$  as:
- $p.d = 0$
  - $p.\text{parent} = p$




# Computing the Distance Map

- We expand each initialized grid cell  $p$ , *i.e.* we take the 8 neighbors.
- For each of these expanded cells  $c_i$

```
dp,c = distance(p, ci)  
if (ci.isEmpty()) {  
    ci.d = dp,c  
    ci.parent = p.parent  
} else if (ci.d > dp,c) {  
    ci.d = dp,c  
    ci.parent = p.parent  
}
```

- Continue iteratively for all the expanded cells.





# Distance Map: Complexity

Computing the distance map requires a time

$$O(\text{grid\_size} * \log(\text{grid\_size}))$$

One query on the distance map requires

$$O(1)$$

Good when

- we have many low dimensional points in the collection
- the grid is small
- we can tolerate small association errors

# Distance Map: Heuristics

Gating:


- expand up to a maximum distance

Best friends:

- construct a distance map also for the measurements and cross check

Lonely Best Friends:

- looking up a region of the distance map instead of a single point



# KD-Trees

What if the points are K-dimensional, with K large?

- The distance map does not scale well:  
*memory grows with K, and so the time to construct it*

KD-Tree:

- search structure that partitions the database according to their spatial distribution
- If the tree is balanced, a search takes  $\log(N)$  with N the number of points

# KD-Trees

Constructing KD-trees can be done with a trivial recursion.

At each time, the set is split in two parts until the number of points in a leaf is smaller than a threshold.

Question:


how to split?

# KD-Trees: simple splitting

Consider the  $k$ -dimensional points as if they were normally distributed.

Compute the **covariance** matrix of the distribution.

Chose a splitting (hyper) plane that passes through the mean and has a normal aligned with the longest axis of the **covariance** matrix.



# KD-Trees

A split plane is characterized by

- a *normal*  $\mathbf{n}_i$
- a *mean*  $\mu_i$

We can check if a point lies on the one side of the plane by evaluating

$$\mathbf{n}_i^T (\mathbf{p}_q - \mu_i) > 0$$

# KD-Trees

Each intermediate node of the tree contains

- The normal of the splitting plane
- The mean
- Pointers to the children nodes

A query requires traversing the tree from the top to the bottom and at each time going left or right.

The query result is an approximation, *i.e.* the tree is a heuristic that might return not the real minimum, depending on how the tree was built.

*Efficient randomized variants (see ANN C++ library)*

# Other Tools

## Projection to lower dimension

- Reduce the dimension of the query points, by projecting them for instance in 2D
- Use some easy heuristic to perform the association in the lower dimensional space

## Bag of words

- Used to determine the appearance similarity of multiple points.
- An item of the search collection contains many points.



# How to speed up the Localizer?

- What would you use?
- Why?
- How?
- What do you expect?

# What about SLAM?

- What would you use?
- Why?
- How?
- What do you expect?